

CS186 Discussion #11

(Transactions & Concurrency Control)

Transactions

- Sequence of reads and writes:
 - `bank_account := bank_account - 500`
 - `wallet := wallet + 500`
- Either all instructions execute or none
- No other instruction sees only part of this execution

ACID

- Atomicity: either none or all instructions executed
- Consistency: database remains in consistent state afterwards
- Isolation: runs as if it is only transaction
- Durability: committed changes are never lost

Isolation

- Could just execute 1 transaction at a time
 - “Serial schedule”
 - Slow
- Maximize parallelism while maintaining sense of isolation
 - Concurrency control

Serializability

- Serial schedule: run 1 transaction at a time
- Equivalent schedules: schedules with same transactions, same final state
- Serializable schedule: schedule equivalent to a serial schedule

T1	wallet1 := wallet1 - 200	acct := acct + 200		
T2			acct := acct - 500	wallet2 := wallet2 + 500

T1	wallet1 := wallet1 - 200	acct := acct + 200		
T2			acct := acct - 500	wallet2 := wallet2 + 500

Serial

T1	wallet1 := wallet1 - 200			acct := acct + 200
T2		acct := acct - 500	wallet2 := wallet2 + 500	

T1	wallet1 := wallet1 - 200			acct := acct + 200
T2		acct := acct - 500	wallet2 := wallet2 + 500	

Serializable

T1	wallet1 := wallet1 - 200			acct := acct + 200
T2		acct := acct - 500	wallet2 := wallet2 + 500	



T1	R(w1)	W(w1)					R(acct)	W(acct)
T2			R(acct)	W(acct)	R(w2)	W(w2)		


T1	R(w1)	W(w1)	R(acct)					W(acct)
T2				R(acct)	W(acct)	R(w2)	W(w2)	

T1	R(w1)	W(w1)	R(acct)					W(acct)
T2				R(acct)	W(acct)	R(w2)	W(w2)	

Not serializable

Conflicts



- Two operations in different transactions on the same object where at least one is a write



3 conflicts

Conflict Serializable


- Conflict serializable: equivalent to some serial schedule with the conflicts in the same order.



Dependency Graph

- One node per transaction
- Edge from T_i to T_j if operation O_i conflicts with O_j and O_i appears earlier in the schedule

T1	R(1)	W(1)	R(a)					W(a)
T2				R(a)	W(a)	R(2)	W(2)	




Dependency Graph

- Theorem: Schedule is conflict serializable iff its dependency graph is acyclic



T1	R(1)	W(1)	R(a)					W(a)
T2				R(a)	W(a)	R(2)	W(2)	

The table illustrates a schedule for two transactions, T1 and T2, across a sequence of operations. T1 performs R(1), W(1), R(a), and W(a). T2 performs R(a), W(a), R(2), and W(2). The dependency graph shows two edges: one from T1's R(a) to T2's R(a), and another from T1's W(a) to T2's W(a).



Dependency Graph

- Theorem: Schedule is conflict serializable iff its dependency graph is acyclic




Serializability vs. Conflict Serializability


- It is much harder to verify serializability
- Any conflict serializable schedule is serializable
- Some serializable schedules are not conflict serializable

Worksheet Page 1


T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)

T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)





Conflict serializable



Conflict serializable


Equivalent serial schedules:


T3, T1, T2

T1, T3, T2

T1	R(A)		R(B)				W(A)	
T2		R(A)		R(B)				W(B)
T3					R(A)			
T4						R(B)		

T1	R(A)		R(B)				W(A)	
T2		R(A)		R(B)				W(B)
T3					R(A)			
T4						R(B)		





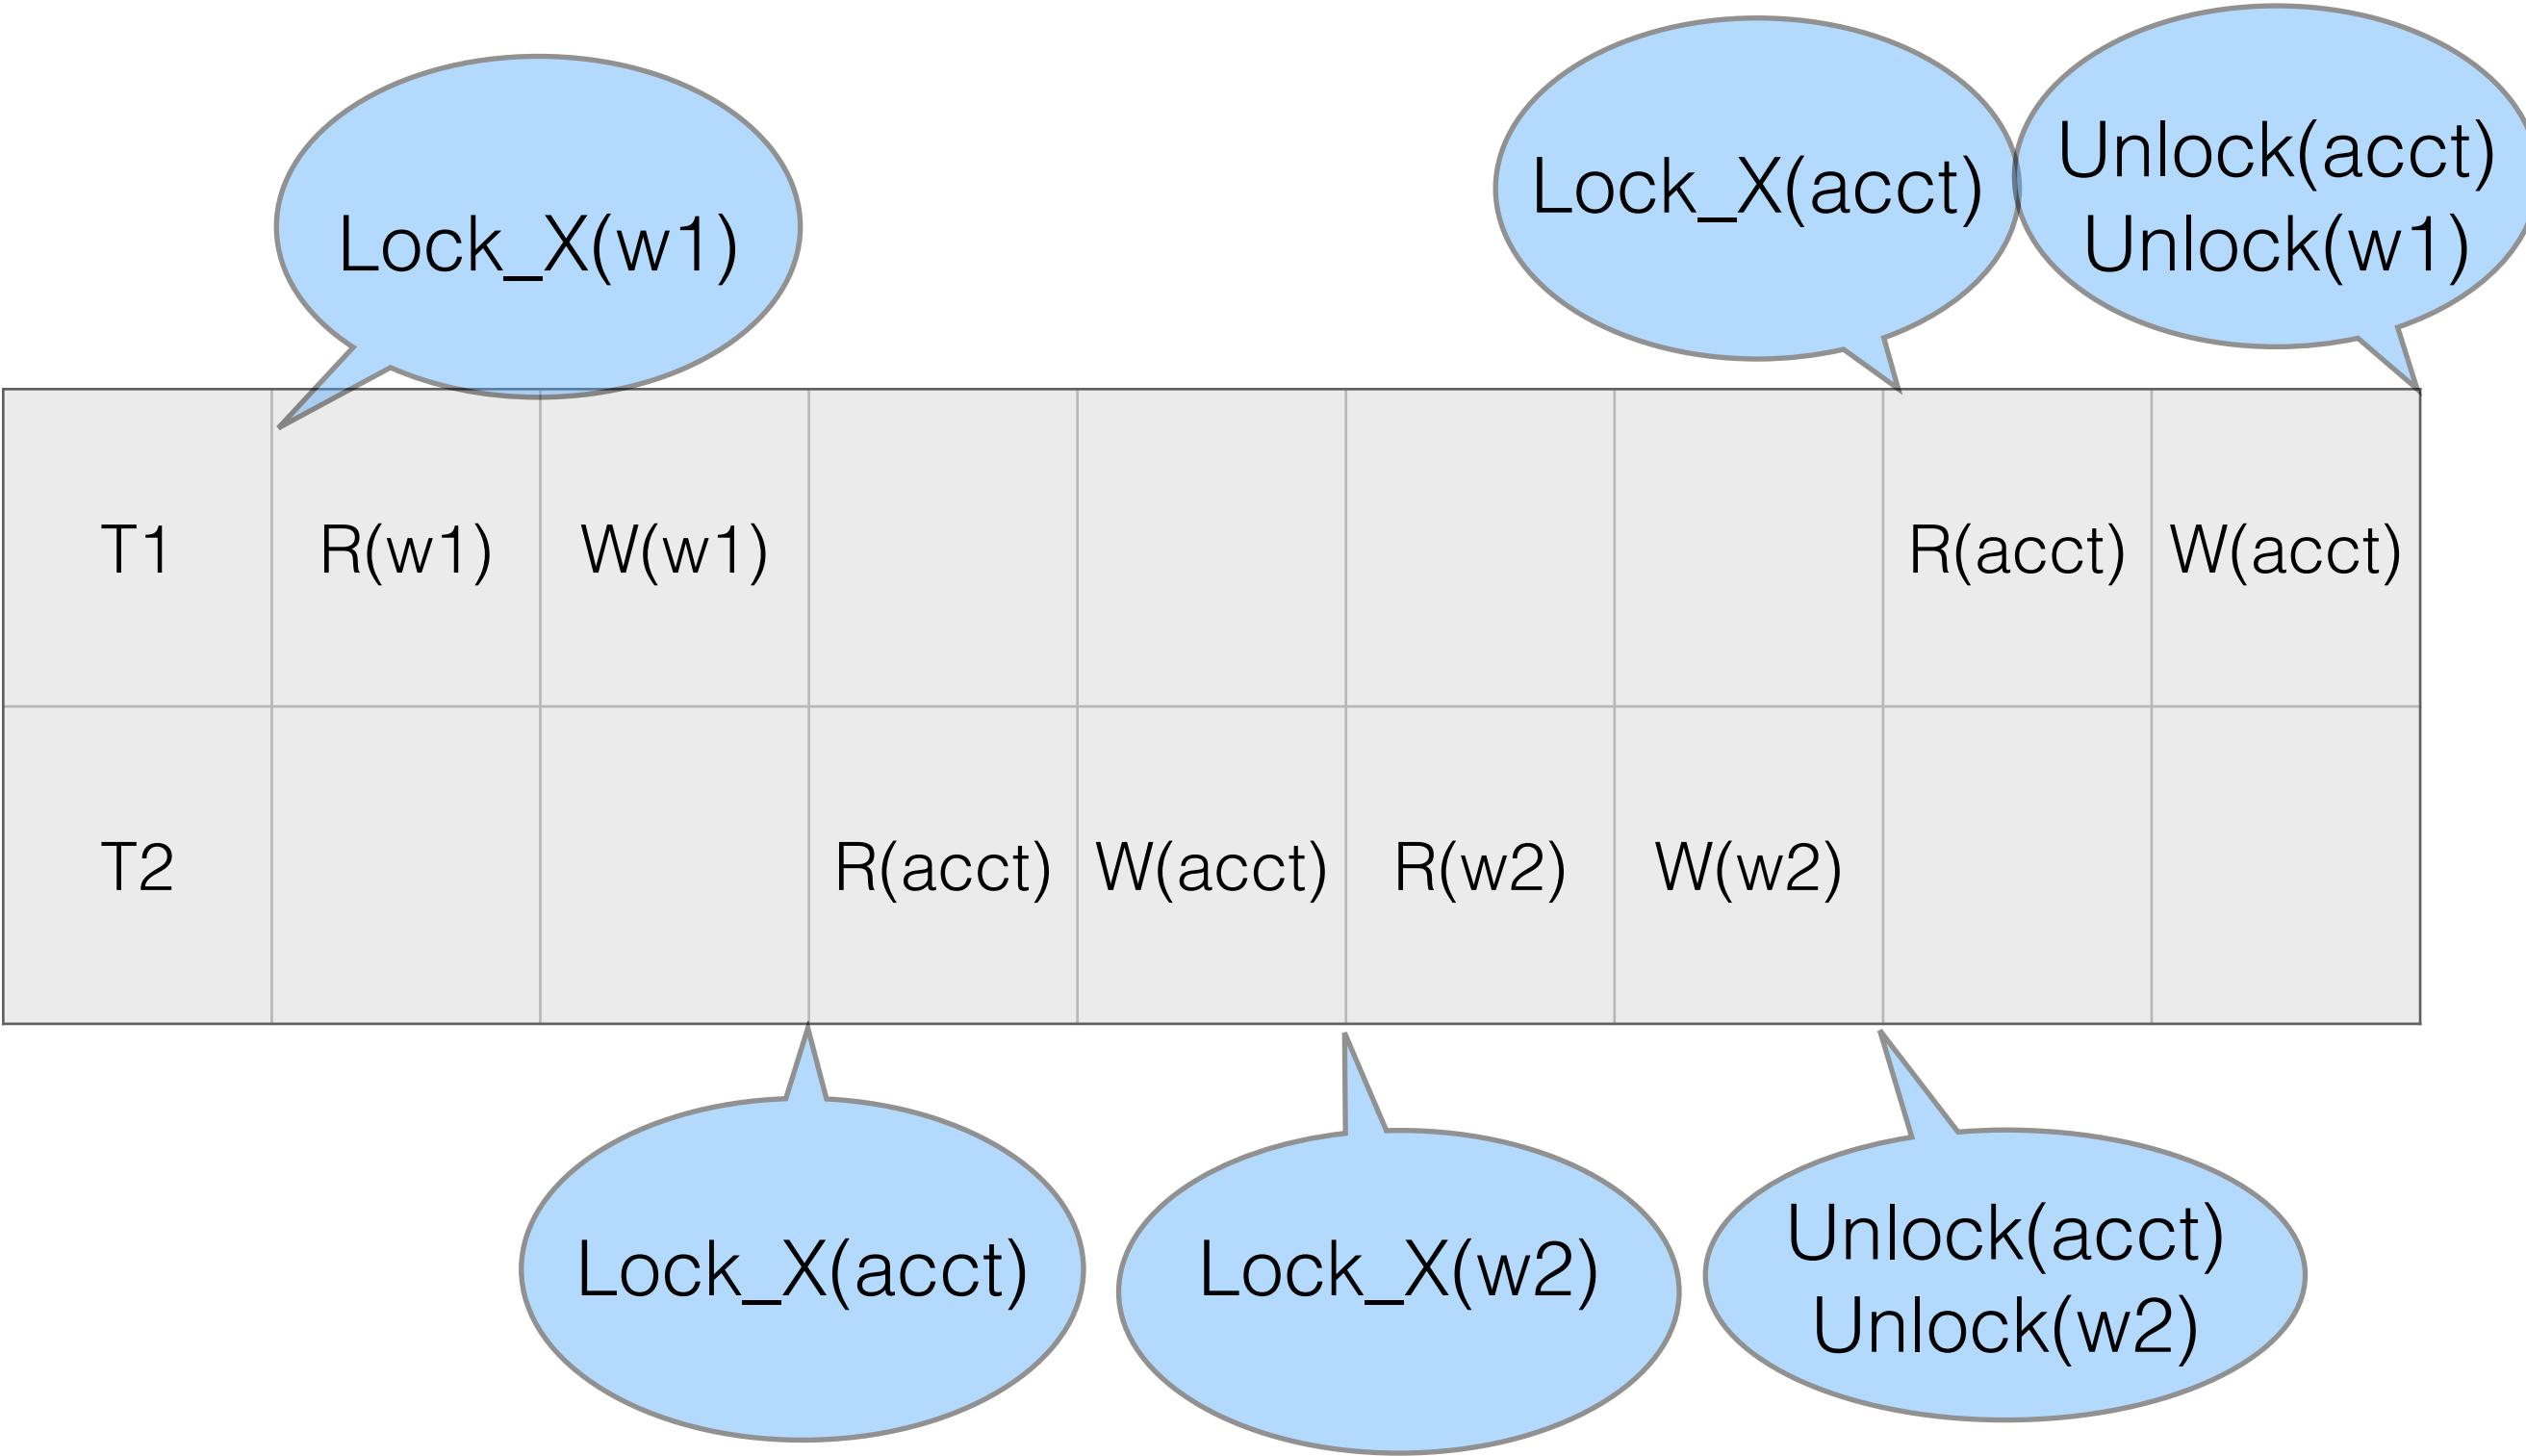
Not conflict serializable,
because cycle in precedence graph.

Locks

- Use locks to control access to objects
- Shared lock: multiple transactions can have shared lock on same item (e.g., reading)
- Exclusive lock: only one lock on item (e.g., writing)

Two-Phase Locking (2PL)

- Rule: Once you release a lock, you may never acquire a new lock.
- Ensures conflict serializability



Strict 2PL

Cascading abort

T1	R(A)	W(A)			Abort
T2			R(A)	W(A)	

Strict two-phase locking: All locks held by transaction are only released at the end of the transaction.

Strict 2PL

Cascading abort

T1	R(A)	W(A)	Abort		
T2				R(A)	W(A)

Strict two-phase locking: All locks held by transaction are only released at the end of the transaction.

Deadlocks

- Dealing with deadlocks:
 - Prevention - stop from occurring
 - Detection - stop while occurring
 - In practice: Timeouts

Time	Transaction 1	Transaction 2
1	Lock_X(A) (granted)	
2		Lock_X(B) (granted)
3	Lock_X(B) (waiting)	
4		Lock_X(A) (waiting)
5

Deadlock Prevention

- Disallow deadlocks from ever occurring
- Use timestamps of transactions to determine which to abort
- If you restart, you restart with original timestamp

Deadlock Prevention

- Two transactions T_{old} and T_{young}
- Wait-Die
 - If T_{old} waiting for a lock from T_{young} , just waits
 - If T_{young} waiting for lock from T_{old} , aborts himself
- Wound-Wait
 - If T_{old} waiting for a lock from T_{young} , aborts T_{young}
 - If T_{young} waiting for lock from T_{old} , just waits

Deadlock Detection


- Create waits-for graph of all transactions
- If cycle exists, abort one of transactions in cycle

Deadlock Detection

T1	X(A)			S(B)			
T2		X(B)	S(A)				
T3					X(C)	S(A)	
T4							S(C)

Deadlock Detection

T1	X(A)			S(B)			
T2		X(B)	S(A)				
T3					X(C)	S(A)	
T4							S(C)



Worksheet Page 2

(B=3, F=300)

Lock_X(B)	
Read(B)	
B = B*10	
Write(B)	Lock_S(B)
Lock_X(F)	
F = B*100	
Write(F)	
Unlock(B)	
Unlock(F)	
	Lock_S(F)
	Read(F)
	Read(B)
	Print(F+B)
	Unlock(B)
	Unlock(F)

(B=3, F=300)

Lock_X(B)	
Read(B)	
B = B*10	
Write(B) B = 30	Lock_S(B)
Lock_X(F)	
F = B*100	
Write(F) F = 3000	
Unlock(B)	
Unlock(F)	
	Lock_S(F)
	Read(F)
	Read(B)
	Print(F+B)
	Unlock(B)
	Unlock(F)

Prints 3030

(B=3, F=300)


Lock_X(B)	
Read(B)	
B = B*10	
Write(B)	Lock_S(B)
Lock_X(F)	
F = B*100	
Write(F)	
Unlock(B)	
Unlock(F)	
	Lock_S(F)
	Read(F)
	Read(B)
	Print(F+B)
	Unlock(B)
	Unlock(F)

Uses Strict 2PL:


Does not release locks until end of transaction

T1	S(A), S(D)		S(B)				
T2		X(B)			X(C)		
T3				S(D), S(C)			X(A)
T4						X(B)	

T1	S(A), S(D)		S(B)				
T2		X(B)			X(C)		
T3				S(D), S(C)			X(A)
T4						X(B)	



T1	S(A), S(D)		S(B)				
T2		X(B)			X(C)		
T3				S(D), S(C)			X(A)
T4						X(B)	



Deadlock
possible:
cycle in
graph